

```

# Filepath
file_path = 'laptop prices.csv'

# Import necessary libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error

# Step 1: Load the dataset
df = pd.read_csv(file_path)

# Data Exploration
print("Dataset Head:\n", df.head())
print("\nDataset Info:\n", df.info())
print("\nSummary Statistics:\n", df.describe())

```

Dataset Head:

	brand	processor_brand	processor_name	processor_gnrtn	ram_gb
0	ASUS	Intel	Core i3	10th	4 GB
1	Lenovo	Intel	Core i3	10th	4 GB
2	Lenovo	Intel	Core i3	10th	4 GB
3	ASUS	Intel	Core i5	10th	8 GB
4	ASUS	Intel	Celeron Dual	Not Available	4 GB

	ssd	hdd	os	os_bit	graphic_card_gb	weight
0	0 GB	1024 GB	Windows	64-bit	0 GB	Casual
1	0 GB	1024 GB	Windows	64-bit	0 GB	Casual
2	0 GB	1024 GB	Windows	64-bit	0 GB	Casual
3	512 GB	0 GB	Windows	32-bit	2 GB	Casual
4	0 GB	512 GB	Windows	64-bit	0 GB	Casual

	Touchscreen	msoffice	Price	rating	Number of Ratings	Number of Reviews
0	No	No	34649	2 stars	3	
0						
1	No	No	38999	3 stars	65	
5						
2	No	No	39999	3 stars	8	
1						
3	No	No	69990	3 stars	0	
0						
4	No	No	26990	3 stars	0	
0						

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 823 entries, 0 to 822

Data columns (total 19 columns):

#	Column	Non-Null Count	Dtype
0	brand	823 non-null	object
1	processor_brand	823 non-null	object
2	processor_name	823 non-null	object
3	processor_gnrtn	823 non-null	object
4	ram_gb	823 non-null	object
5	ram_type	823 non-null	object
6	ssd	823 non-null	object
7	hdd	823 non-null	object
8	os	823 non-null	object
9	os_bit	823 non-null	object
10	graphic_card_gb	823 non-null	object
11	weight	823 non-null	object
12	warranty	823 non-null	object
13	Touchscreen	823 non-null	object
14	msoffice	823 non-null	object
15	Price	823 non-null	int64
16	rating	823 non-null	object
17	Number of Ratings	823 non-null	int64
18	Number of Reviews	823 non-null	int64

dtypes: int64(3), object(16)

memory usage: 122.3+ KB

Dataset Info:

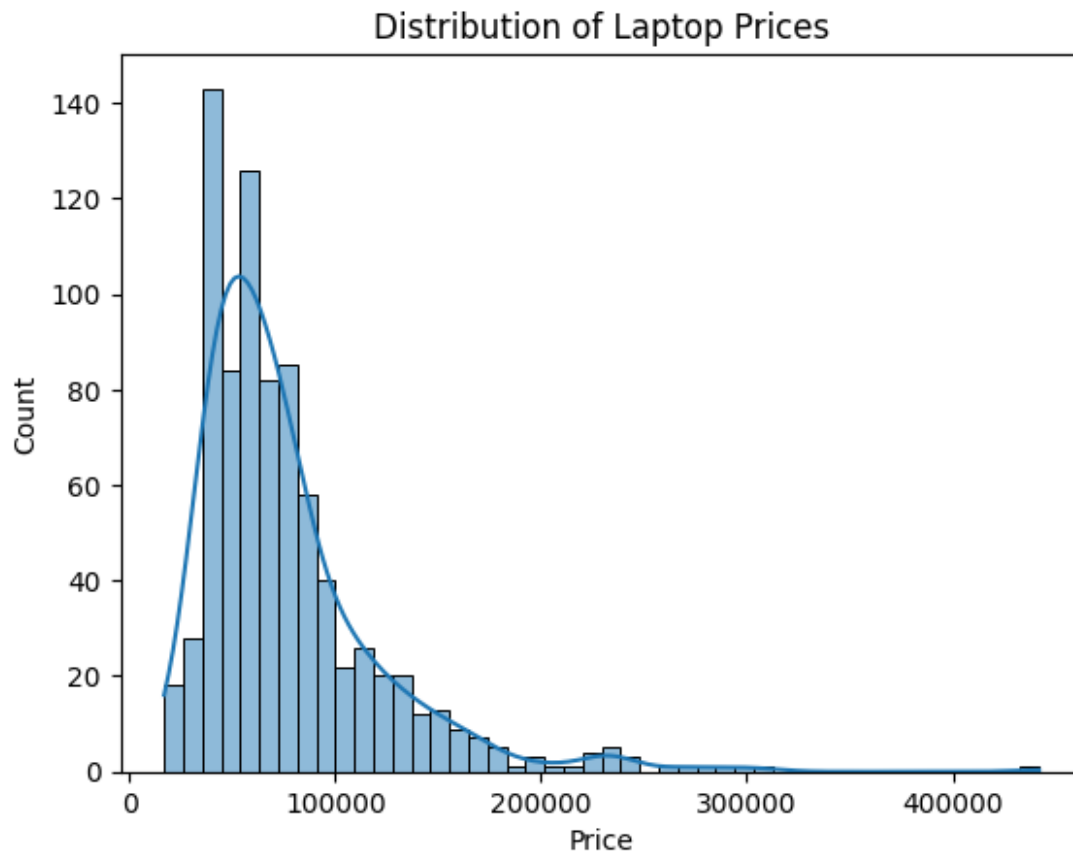
None

Summary Statistics:

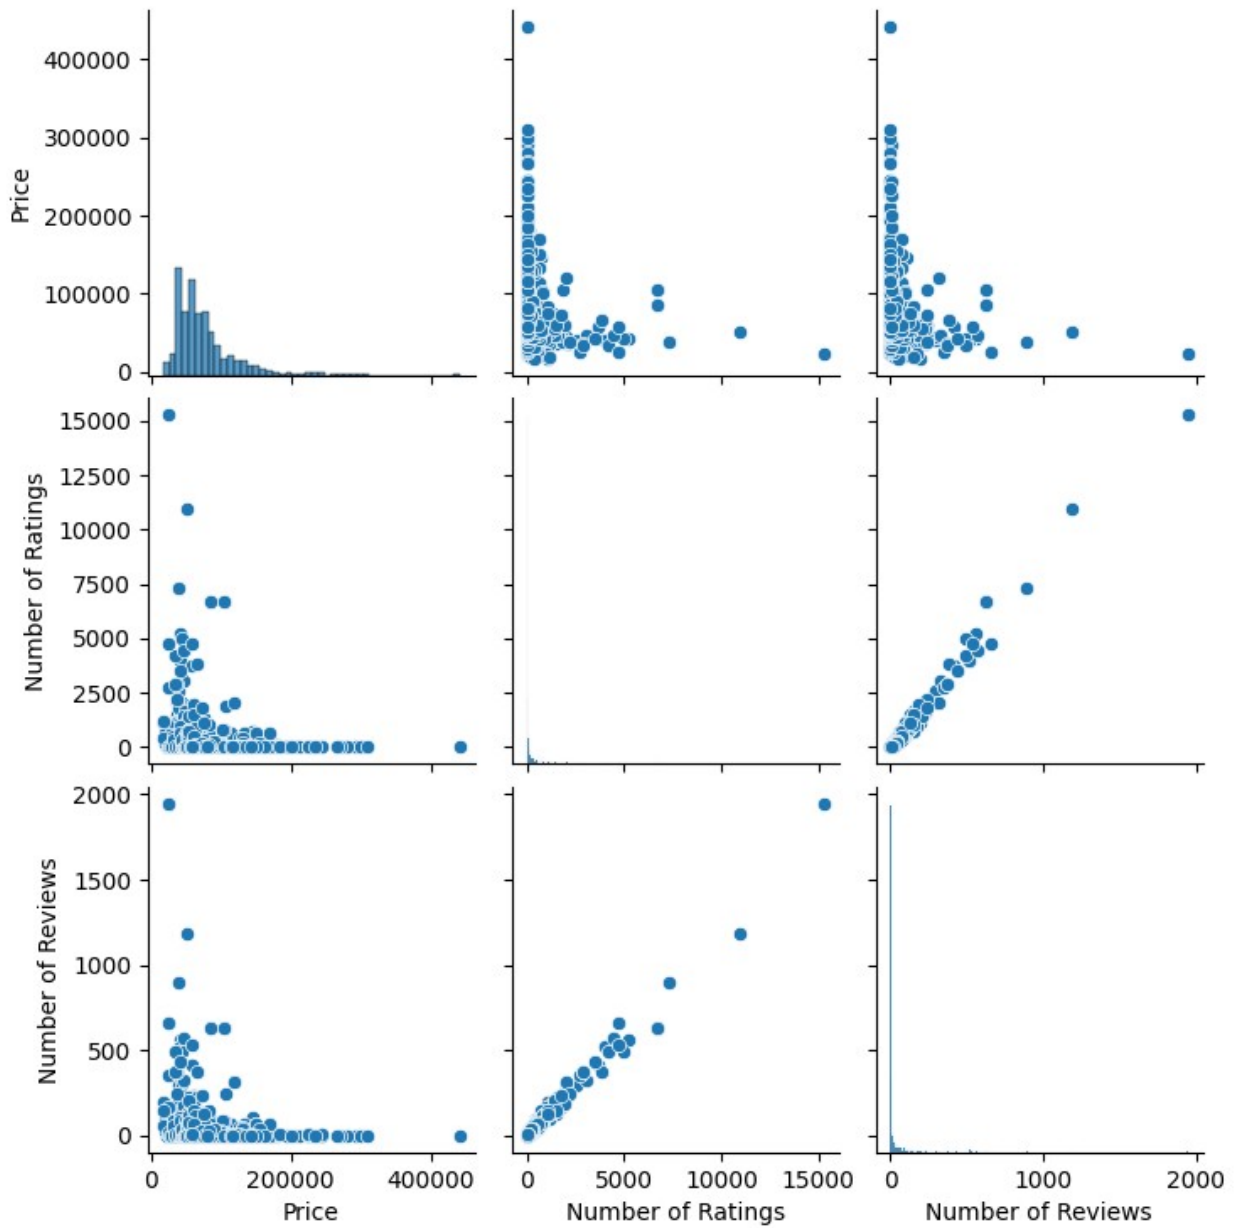
	Price	Number of Ratings	Number of Reviews
count	823.000000	823.000000	823.000000
mean	76745.177400	315.301337	37.609964
std	45101.790525	1047.382654	121.728017
min	16990.000000	0.000000	0.000000
25%	46095.000000	0.000000	0.000000

50%	64990.000000	17.000000	2.000000
75%	89636.000000	139.500000	18.000000
max	441990.000000	15279.000000	1947.000000

```
# Visualize the distribution of the target variable (Price)
sns.histplot(df['Price'], kde=True)
plt.title('Distribution of Laptop Prices')
plt.show()
```



```
# Visualize relationships between features and target variable
sns.pairplot(df)
plt.show()
```



```
# Step 2: Data Preprocessing and Feature Engineering
# Check for missing values
print("\nMissing Values:\n", df.isnull().sum())
```

```
Missing Values:
brand          0
processor_brand 0
processor_name  0
processor_gnrtn 0
ram_gb         0
ram_type       0
```

```

ssd                0
hdd                0
os                 0
os_bit             0
graphic_card_gb    0
weight             0
warranty           0
Touchscreen        0
msoffice           0
Price              0
rating             0
Number of Ratings  0
Number of Reviews  0
dtype: int64

# Handle missing values (if any)
df = df.dropna() # Dropping rows with missing values for simplicity

# Convert categorical columns to numerical
label_encoders = {}
for column in df.select_dtypes(include=['object']).columns:
    le = LabelEncoder()
    df[column] = le.fit_transform(df[column])
    label_encoders[column] = le

# Feature scaling
scaler = StandardScaler()
df[df.columns] = scaler.fit_transform(df[df.columns])

# Step 3: Split the data into training and testing sets
X = df.drop('Price', axis=1) # Features
y = df['Price'] # Target variable

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Step 4: Model Selection and Training
models = {
    'Linear Regression': LinearRegression(),
    'Random Forest': RandomForestRegressor(n_estimators=100,
random_state=42),
    'Gradient Boosting': GradientBoostingRegressor(n_estimators=100,
random_state=42)
}

trained_models = {}
for model_name, model in models.items():
    model.fit(X_train, y_train)
    trained_models[model_name] = model

```

### *# Step 5: Model Evaluation*

```
def evaluate_model(model, X_test, y_test):  
    y_pred = model.predict(X_test)  
    mae = mean_absolute_error(y_test, y_pred)  
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))  
    return mae, rmse
```

```
for model_name, model in trained_models.items():  
    mae, rmse = evaluate_model(model, X_test, y_test)  
    print(f"{model_name} - MAE: {mae}, RMSE: {rmse}")
```

```
Linear Regression - MAE: 0.432176348476466, RMSE: 0.6885798734737727  
Random Forest - MAE: 0.2860755773765587, RMSE: 0.5846661440561662  
Gradient Boosting - MAE: 0.31462902291940364, RMSE: 0.601076982247702
```

### *# Cross-validation*

```
def cross_val(model, X, y):  
    scores = cross_val_score(model, X, y,  
        scoring='neg_mean_squared_error', cv=5)  
    rmse_scores = np.sqrt(-scores)  
    return rmse_scores
```

```
for model_name, model in models.items():  
    scores = cross_val(model, X, y)  
    print(f"{model_name} - Cross-Validation RMSE: {scores.mean()}")
```

```
Linear Regression - Cross-Validation RMSE: 0.7828636881582345  
Random Forest - Cross-Validation RMSE: 0.7368779115436246  
Gradient Boosting - Cross-Validation RMSE: 0.6880702219334448
```

### *# Sample prediction*

```
sample_data = X_test.iloc[0].values.reshape(1, -1)  
predicted_price = trained_models['Linear  
Regression'].predict(sample_data)  
print(f'Sample Predicted Price: {predicted_price[0]}')
```

```
Sample Predicted Price: 0.6900953951677803
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:  
UserWarning: X does not have valid feature names, but LinearRegression  
was fitted with feature names  
    warnings.warn(  

```

### *# Save the trained models, encoders, and scaler*

```
import pickle
```

```
for model_name, model in trained_models.items():  
    with open(f'/content/{model_name.replace(" ",  
        "_").lower()}_model.pkl', 'wb') as file:  
        pickle.dump(model, file)
```

```
with open('/content/label_encoders.pkl', 'wb') as file:  
    pickle.dump(label_encoders, file)  
  
with open('/content/scaler.pkl', 'wb') as file:  
    pickle.dump(scaler, file)
```