

```

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score
import pandas as pd

```

Step 1: Data Preparation

Load your dataset

```
data = pd.read_csv('emails.csv')
```

```
print(data.columns)
```

```
print(data.head())
```

```
Index(['Email No.', 'the', 'to', 'ect', 'and', 'for', 'of', 'a', 'you', 'hou',
      'you', 'hou',

```

```

      ...
      'connevey', 'jay', 'valued', 'lay', 'infrastructure',
'military',
      'allowing', 'ff', 'dry', 'Prediction'],
      dtype='object', length=3002)

```

	Email No.	the	to	ect	and	for	of	a	you	hou	...	connevey
0	Email 1	0	0	1	0	0	0	2	0	0	...	0
1	Email 2	8	13	24	6	6	2	102	1	27	...	0
2	Email 3	0	0	1	0	0	0	8	0	0	...	0
3	Email 4	0	5	22	0	5	1	51	2	10	...	0
4	Email 5	7	6	17	1	5	2	57	0	9	...	0

	valued	lay	infrastructure	military	allowing	ff	dry
0	0	0	0	0	0	0	0
1	0	0	0	0	0	1	0
2	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0

```

0
4      0      0      0      0      0      1      0
0

[5 rows x 3002 columns]

# Assign labels to y (target variable)
y = data['Prediction']

# Prepare features (X)
X = data.drop(columns=['Email No.', 'Prediction'])

# Handle missing values
imputer = SimpleImputer(strategy='mean')
X_imputed = imputer.fit_transform(X)

# Scale numerical features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_imputed)

# Step 2: Model Selection

# Initialize classifiers
classifiers = {
    'Logistic Regression': LogisticRegression(),
    'Decision Tree': DecisionTreeClassifier(),
    'Support Vector Machine': SVC(),
    'Random Forest': RandomForestClassifier(),
    'Gradient Boosting': GradientBoostingClassifier()
}

# Step 3: Model Training and Evaluation

for name, clf in classifiers.items():
    print(f"Training and evaluating {name}...")

    # Split the dataset into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.2, random_state=42)

    # Train the classifier
    clf.fit(X_train, y_train)

Training and evaluating Logistic Regression...

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as

```

shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Training and evaluating Decision Tree...

Training and evaluating Support Vector Machine...

Training and evaluating Random Forest...

Training and evaluating Gradient Boosting...

```
# Model Training and Evaluation
```

```
# Model Training and Evaluation
```

```
for name, clf in classifiers.items():
```

```
    print(f"Training and evaluating {name}...")
```

```
    clf.fit(X_train, y_train)
```

```
    predictions = clf.predict(X_test)
```

```
    # Evaluation Metrics
```

```
    accuracy = accuracy_score(y_test, predictions)
```

```
    precision = precision_score(y_test, predictions)
```

```
    recall = recall_score(y_test, predictions, pos_label=1) #
```

```
    Assuming positive label is 1
```

```
    f1 = f1_score(y_test, predictions, pos_label=1) # Assuming  
    positive label is 1
```

```
    print(f"Accuracy: {accuracy:.2f}")
```

```
    print(f"Precision: {precision:.2f}")
```

```
    print(f"Recall: {recall:.2f}")
```

```
    print(f"F1-Score: {f1:.2f}")
```

```
    print()
```

Training and evaluating Logistic Regression...

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/  
_logistic.py:458: ConvergenceWarning: lbfgs failed to converge  
(status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Accuracy: 0.97
Precision: 0.93
Recall: 0.97
F1-Score: 0.95

Training and evaluating Decision Tree...

Accuracy: 0.93
Precision: 0.88
Recall: 0.87
F1-Score: 0.87

Training and evaluating Support Vector Machine...

Accuracy: 0.95
Precision: 1.00
Recall: 0.82
F1-Score: 0.90

Training and evaluating Random Forest...

Accuracy: 0.98
Precision: 0.95
Recall: 0.97
F1-Score: 0.96

Training and evaluating Gradient Boosting...

Accuracy: 0.97
Precision: 0.94
Recall: 0.96
F1-Score: 0.95

Step 4: Model Evaluation with Cross-Validation

```
for name, clf in classifiers.items():  
    print(f"Evaluating {name} with cross-validation...")  
  
    # Use cross-validation to evaluate the model's generalization  
    # ability  
    cv_scores = cross_val_score(clf, X_scaled, y, cv=5,  
                                scoring='accuracy')  
  
    # Print the mean cross-validation score  
    print(f"Mean Cross-Validation Accuracy: {cv_scores.mean():.2f}")
```

Evaluating Logistic Regression with cross-validation...

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/
_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as

shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic
.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic
.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic
.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic
.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Mean Cross-Validation Accuracy: 0.96

Evaluating Decision Tree with cross-validation...

Mean Cross-Validation Accuracy: 0.91

Evaluating Support Vector Machine with cross-validation...

Mean Cross-Validation Accuracy: 0.92

Evaluating Random Forest with cross-validation...

Mean Cross-Validation Accuracy: 0.96

Evaluating Gradient Boosting with cross-validation...

Mean Cross-Validation Accuracy: 0.95

```
import matplotlib.pyplot as plt
```

```
# Initialize lists to store evaluation metrics
```

```
classifier_names = []
```

```
accuracies = []
```

```
precisions = []
```

```
recalls = []
```

```
f1_scores = []
```

```
# Model Training and Evaluation
```

```
for name, clf in classifiers.items():
```

```
    print(f"Training and evaluating {name}...")
```

```
    clf.fit(X_train, y_train)
```

```
    predictions = clf.predict(X_test)
```

```
# Evaluation Metrics
```

```
    accuracy = accuracy_score(y_test, predictions)
```

```
    precision = precision_score(y_test, predictions)
```

```
    recall = recall_score(y_test, predictions)
```

```
    f1 = f1_score(y_test, predictions)
```

```
# Append metrics to lists
```

```
    classifier_names.append(name)
```

```
    accuracies.append(accuracy)
```

```
    precisions.append(precision)
```

```
    recalls.append(recall)
```

```
    f1_scores.append(f1)
```

```
# Plotting
```

```
plt.figure(figsize=(10, 6))
```

```
# Accuracy
```

```
plt.plot(classifier_names, accuracies, marker='o', label='Accuracy')
```

```

# Precision
plt.plot(classifier_names, precisions, marker='o', label='Precision')

# Recall
plt.plot(classifier_names, recalls, marker='o', label='Recall')

# F1-score
plt.plot(classifier_names, f1_scores, marker='o', label='F1-Score')

# Labels and Title
plt.xlabel('Classifier')
plt.ylabel('Score')
plt.title('Model Evaluation Metrics')
plt.xticks(rotation=45)
plt.legend()
plt.grid(True)

# Show plot
plt.tight_layout()
plt.show()

```

Training and evaluating Logistic Regression...

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Training and evaluating Decision Tree...

Training and evaluating Support Vector Machine...

Training and evaluating Random Forest...

Training and evaluating Gradient Boosting...

